

The First AI Was a Book. The Next One Is Being Locked Away.

We're arguing about AI escape while companies shred the commons that taught us to think.

By Nicholas B. Carty | September 27, 2026

Anthropic bought millions of books, cut off their spines, scanned them, and threw the bodies away. Court filings say it plainly. That is not a sci-fi plot. It is enclosure.

The panic says AI will escape our control. The receipts say something quieter and more dangerous: we are being locked out of the memory AI runs on.

I wrote last week that the first AI was a book. Words were put into books to educate and change thinking. Same function, different century. From pages to interfaces. Same function, new substrate. The danger is not AI going out of control. The danger is not being part of it. Exclusion, not escape.

History backs the claim. Writing let us store information outside the brain for the first time, with readable clay tablets appearing around 3,400 B.C. That was cognitive outsourcing on day one. We could not hold the tax rolls or star charts in our heads, so we pressed them into clay.

Philosophers later called this the Extended Mind. Notebooks are not tools. They are parts of the mind. Books are minds you can hand to someone else.

The substrate kept changing. Papyrus, parchment, paper, tape, hard drives, SSDs that now hold about a trillion bits per gram. The job never changed: hold memory outside the skull so many people can think together across time.

Interfaces are just books that answer back. A search bar is a moving index. A chatbot is a margin note that learned to talk. The shift from page to prompt does not break the lineage. It proves it.

For three years my team has worked the bottleneck where this breaks: drift, memory, domains. In our reference implementation we had to solve a boring, brutal problem. If a system can only move forward, a single bad sensor reading that passes validation will haunt every future answer. We call the symptoms Ghost Invariants and Broken Lineage.

Our fix is not more moderation. It is a spine. Every state change is sealed in a content-addressed Merkle node with a cryptographic root. You can prove the state at step N, detect divergence, and cleanly revert to the last good root. Think of it as a library's revert button, built in code.

That matters because the real risk is not rogue AGI. It is systematic exclusion from the extended mind. The GSMA warns the greatest divide is no longer between countries but

between people who can afford to participate in the digital economy and those who cannot. Billions risk missing the next generation of services before they ever get online.

It is worse than a device gap. New research shows inequality is now about capacity, not just access: who can understand, critique, and meaningfully use AI-mediated systems. You can own a phone and still be locked out.

If the first AI was a book, the first AI safety measure was the public library. We did not make books safe by dulling them. We made them common. Legal deposit. Fair use. Interlibrary loan.

We need the same for interfaces, and it comes down to two moves.

First, memory integrity. Preserve training corpora with cryptographic anchoring, versioned datasets, and public diff logs. If you ingest the commons, you owe the commons a copy. Implement revertible state, so drift is measurable and fixable, not permanent.

Second, domain openness. Treat foundation models used in education, health, law, and public work as civic infrastructure. That means library-grade access: reasonable cost, audit rights, and the ability to fork to a prior signed root when divergence is proven.

This is not anti-tech. It is pro-cognition. The arc from clay to cloud shows that outsourcing memory makes us smarter, but only if the memory stays ours to read, write, and contest.

The debate keeps asking how we control AI. The better question is how we keep people inside it. Build the spine. Keep people in the room.